Using Big Data to Improve Product Quality

Thomas Debeauvais

tdebeauv@uci.edu

Outline

- Netflix recommendations
- What is big data?
- Industry examples
- A technical example: singular value decomposition for recommender systems
- Thought exercises

"Connecting people to the movies they love"

Critically-acclaimed Fight-the-System Documentaries

Based on your interest in...

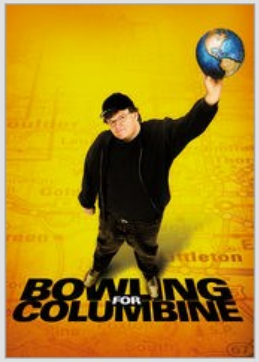

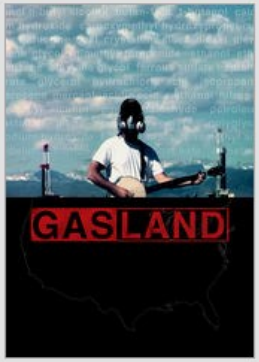

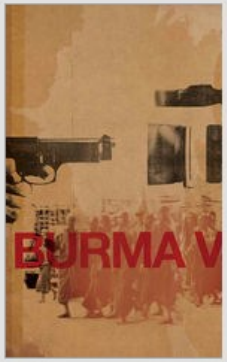

Teen TV Dramas Featuring a Strong Female Lead

Your taste preferences created this row.

TV Dramas
Strong Women.

As well as your interest in...

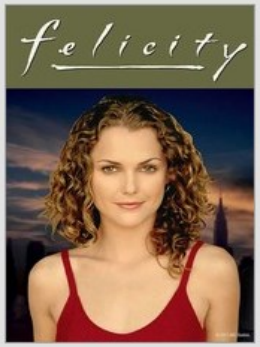

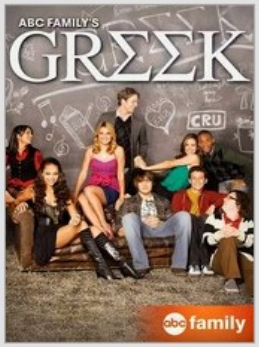

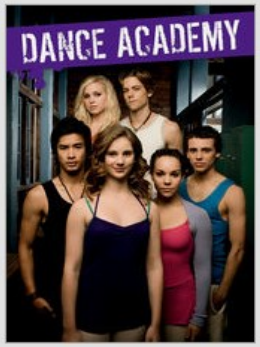

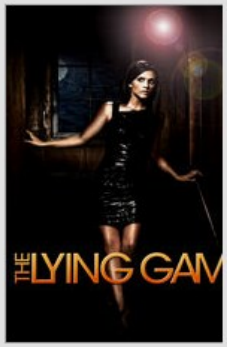

Movies Featuring an Epic Nicolas Cage Meltdown

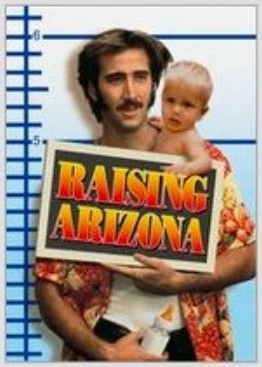

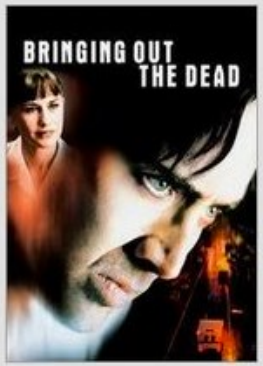

http://genresofnetflix.tumblr.com/

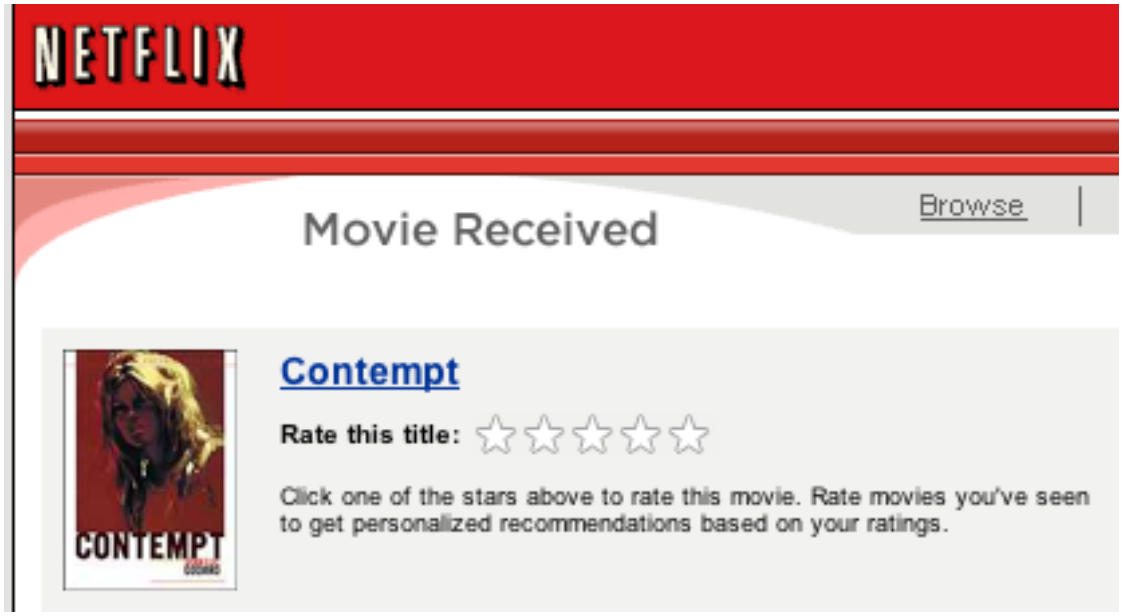
NETFLIX
Movie Received
Browse
Contempt
Rate this title:
Click one of the stars above to rate this movie. Rate movies you've seen to get personalized recommendations based on your ratings.
CONTEMPT

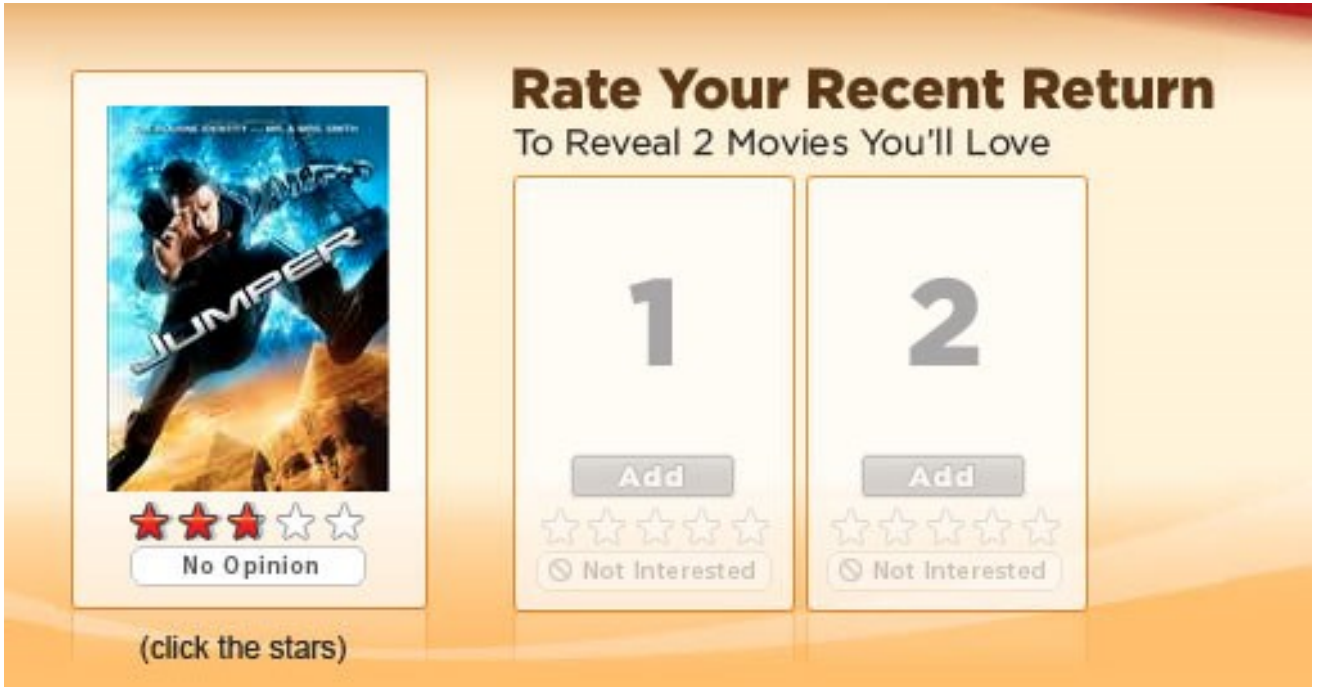
Rate Your Recent Return
To Reveal 2 Movies You'll Love
JUMPER
No Opinion
(click the stars)
1
2
Add
Add
Not Interested
Not Interested

Intuitions

1- People rate movies high

*They **generally** don't bother rating bad movies*

People rate movies high

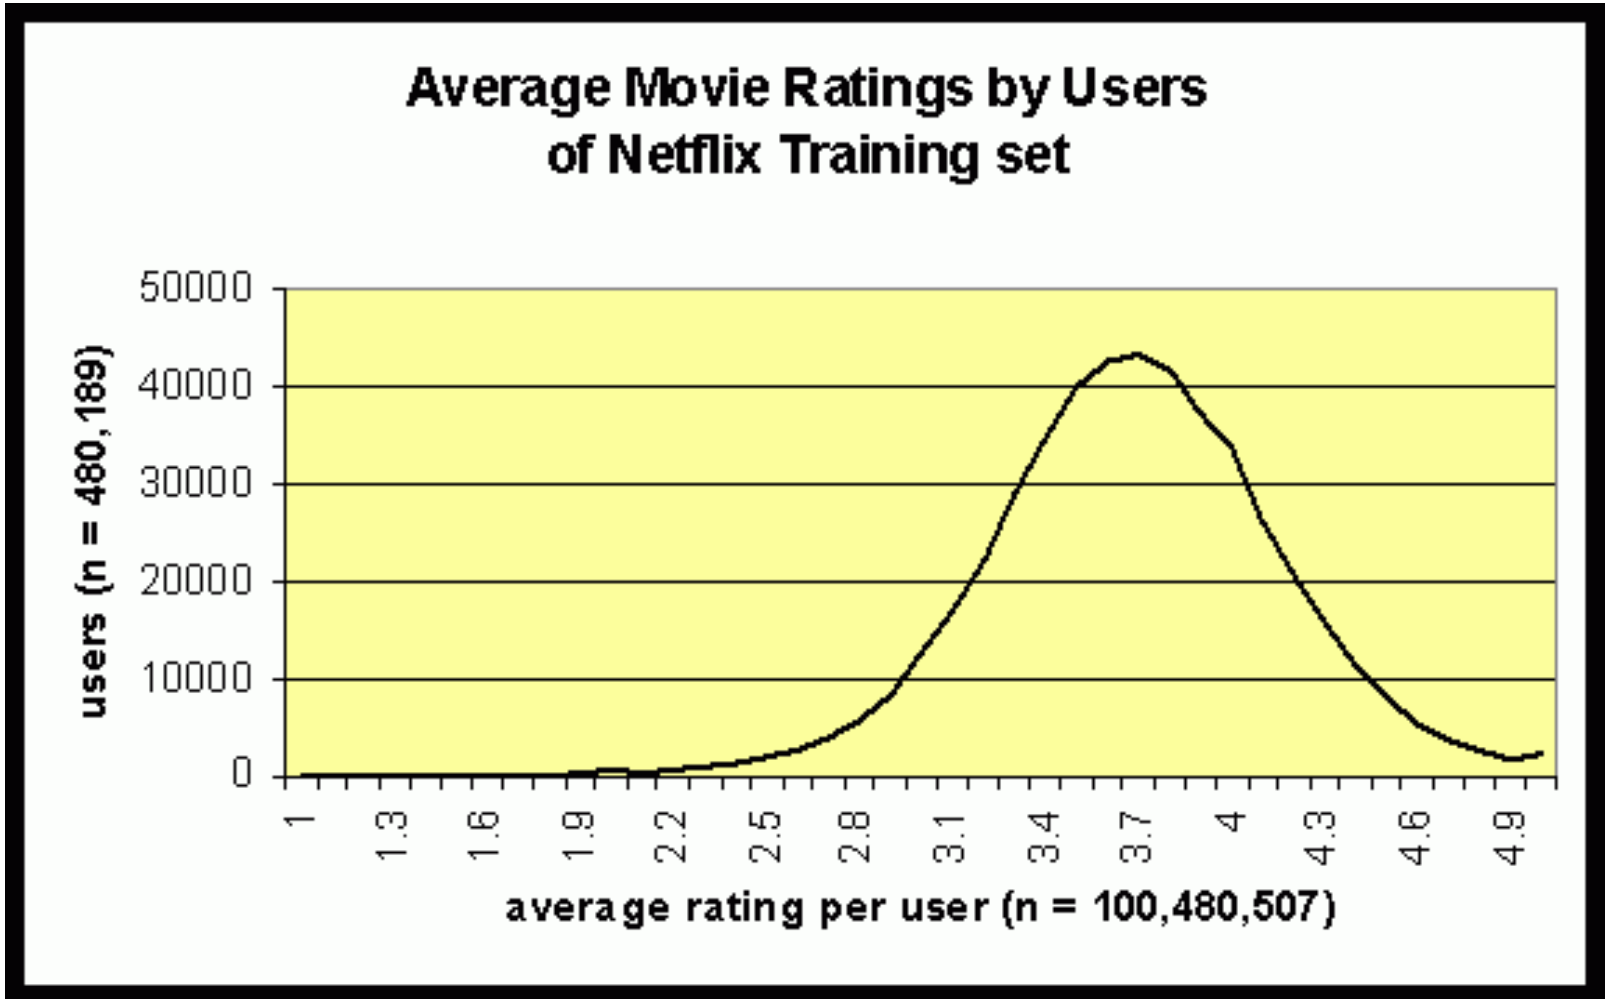

http://www.netflixprize.com/community/viewtopic.php?pid=5941#p5941

Intuitions

2- Older movies are rated higher

Nostalgia

Netflix keeps only the best of old movies

Older movies are rated higher

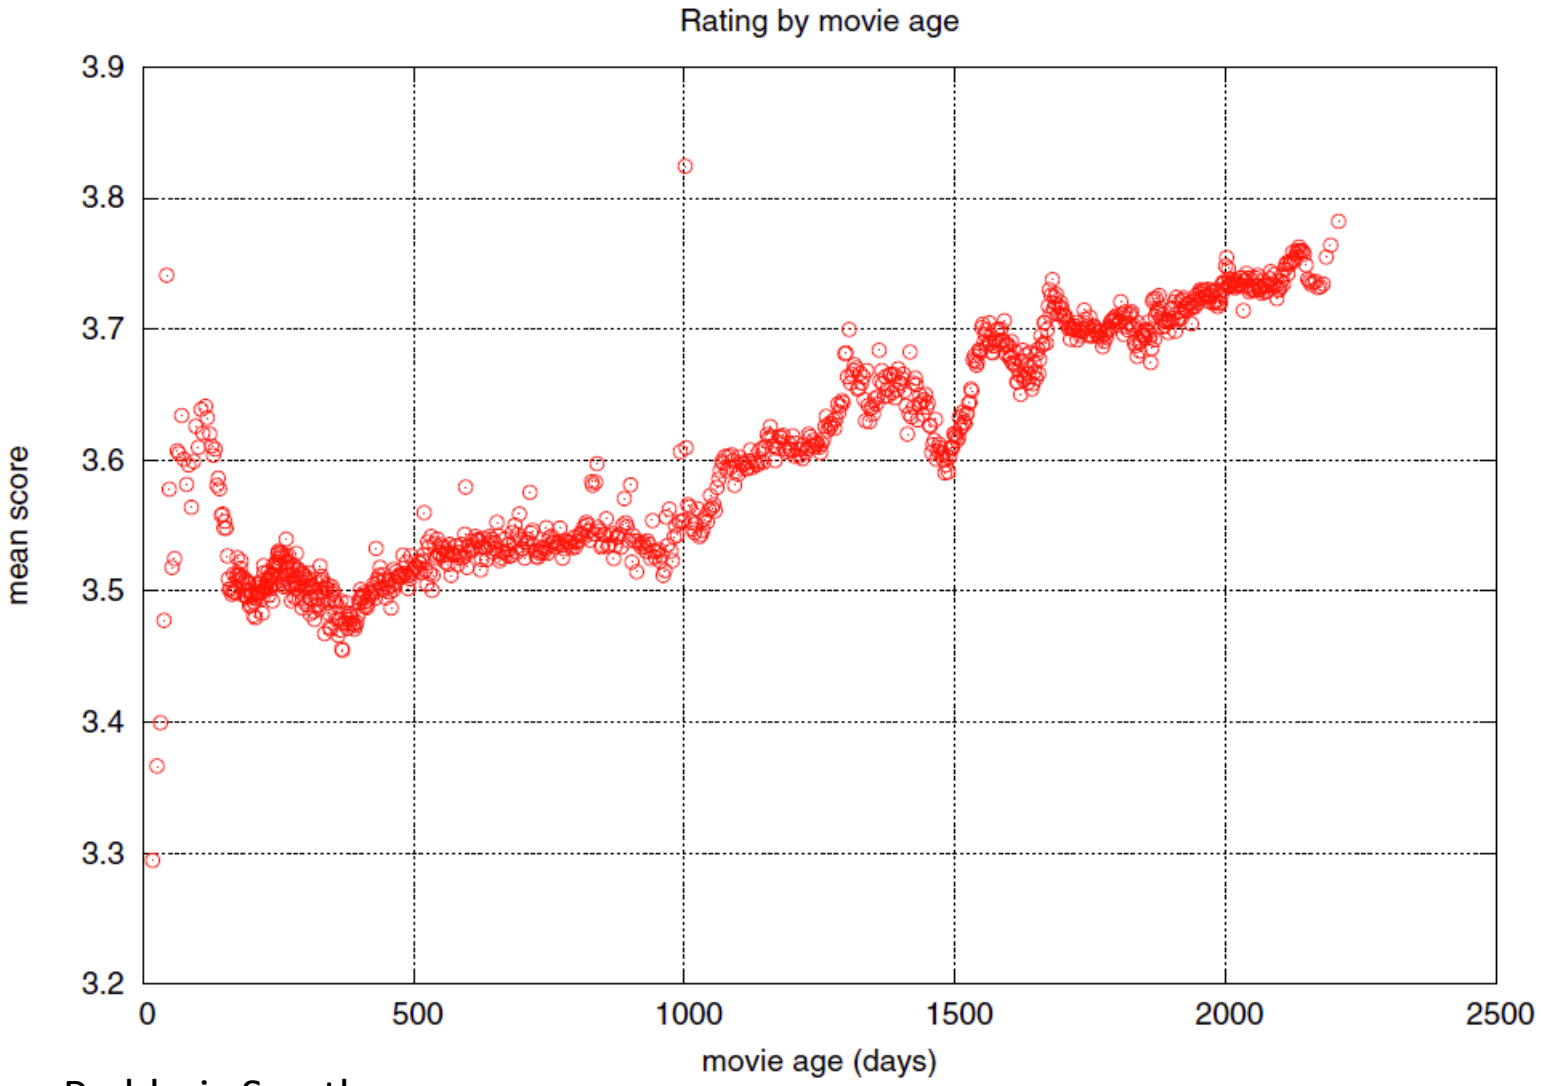

From Paddraic Smyth

Intuitions

3- Movies have genres

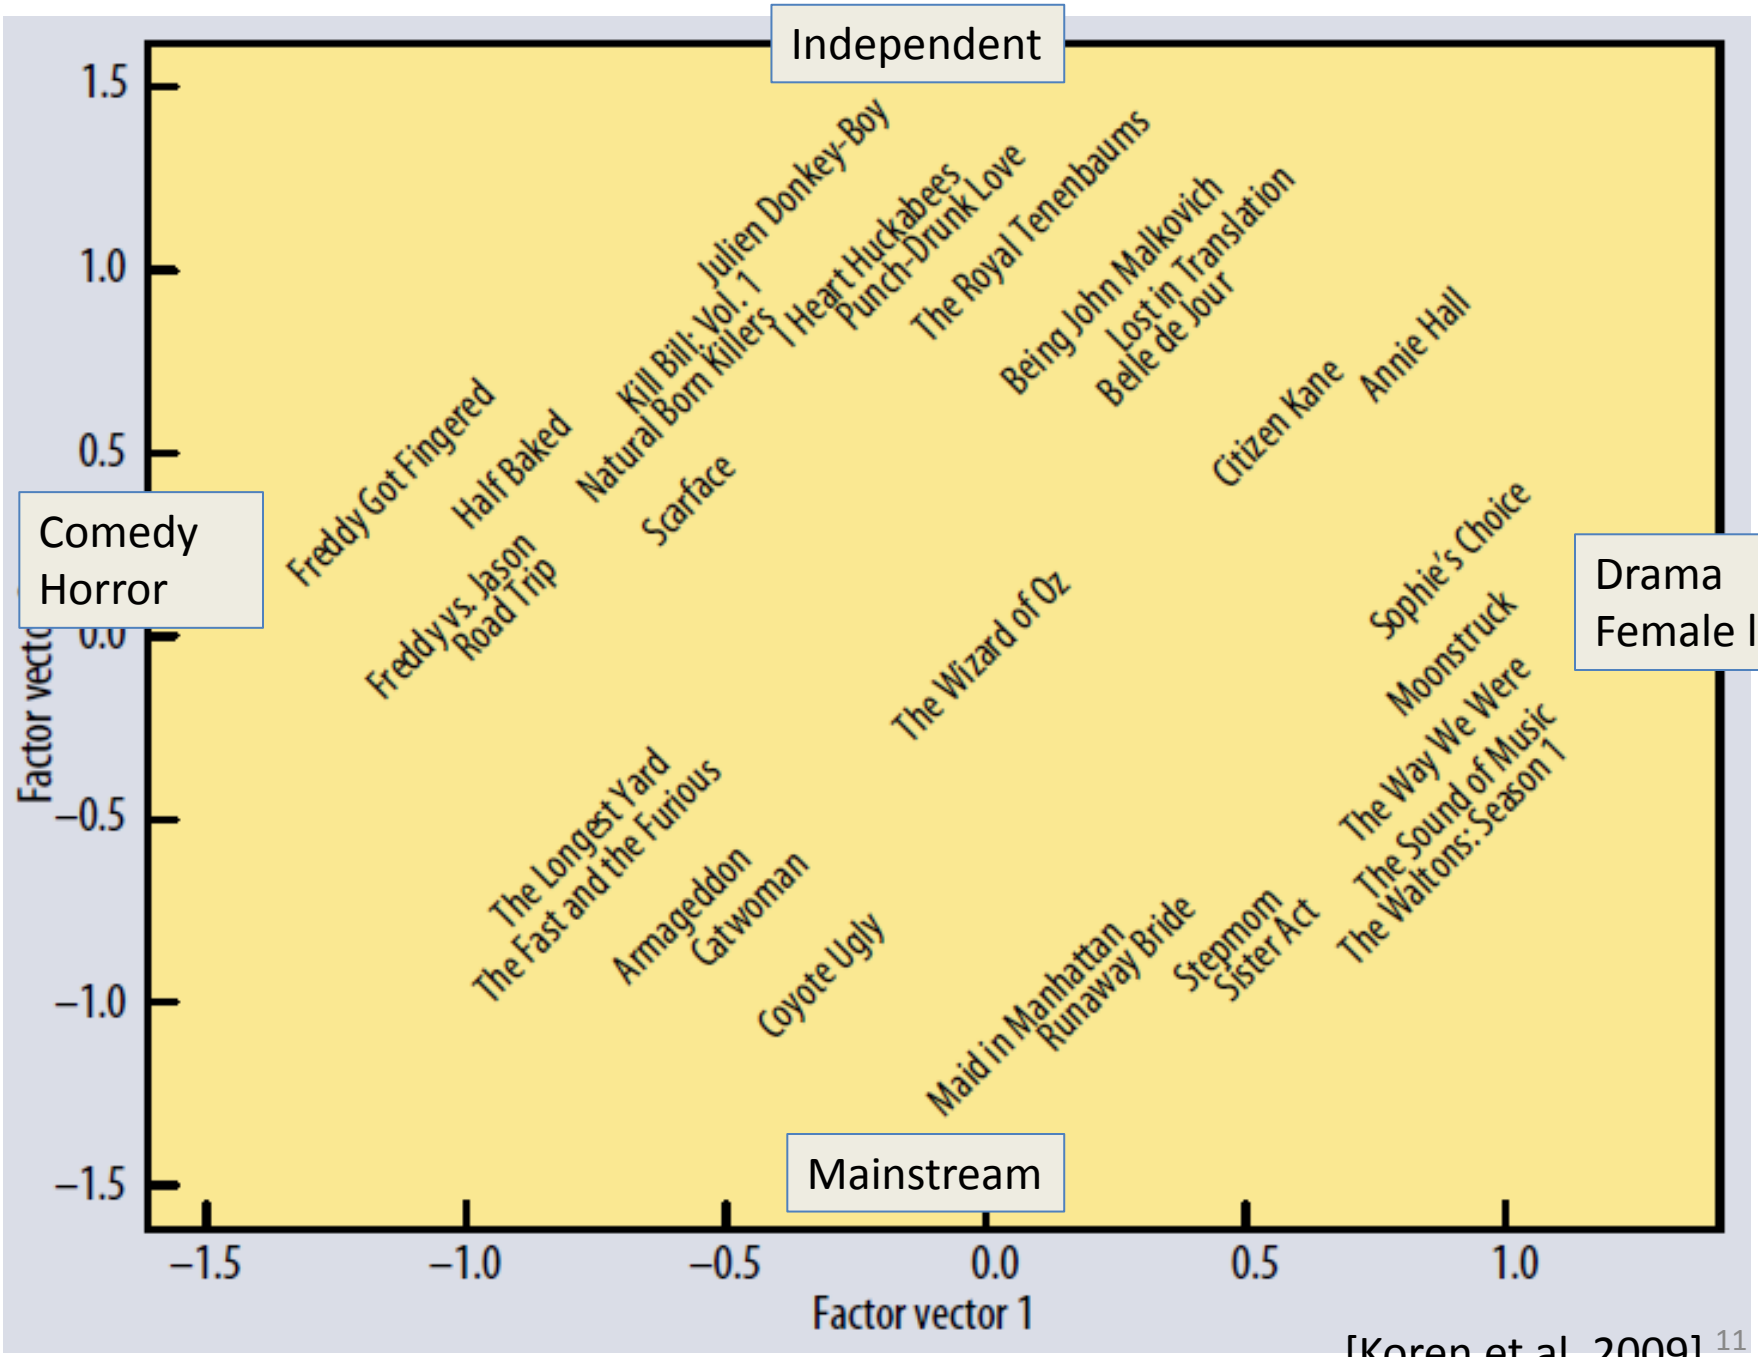

[Koren et al. 2009]

Take-away

- Confirm or reject intuitions with data

WHAT IS BIG DATA?


BIG
DATA

Data mining

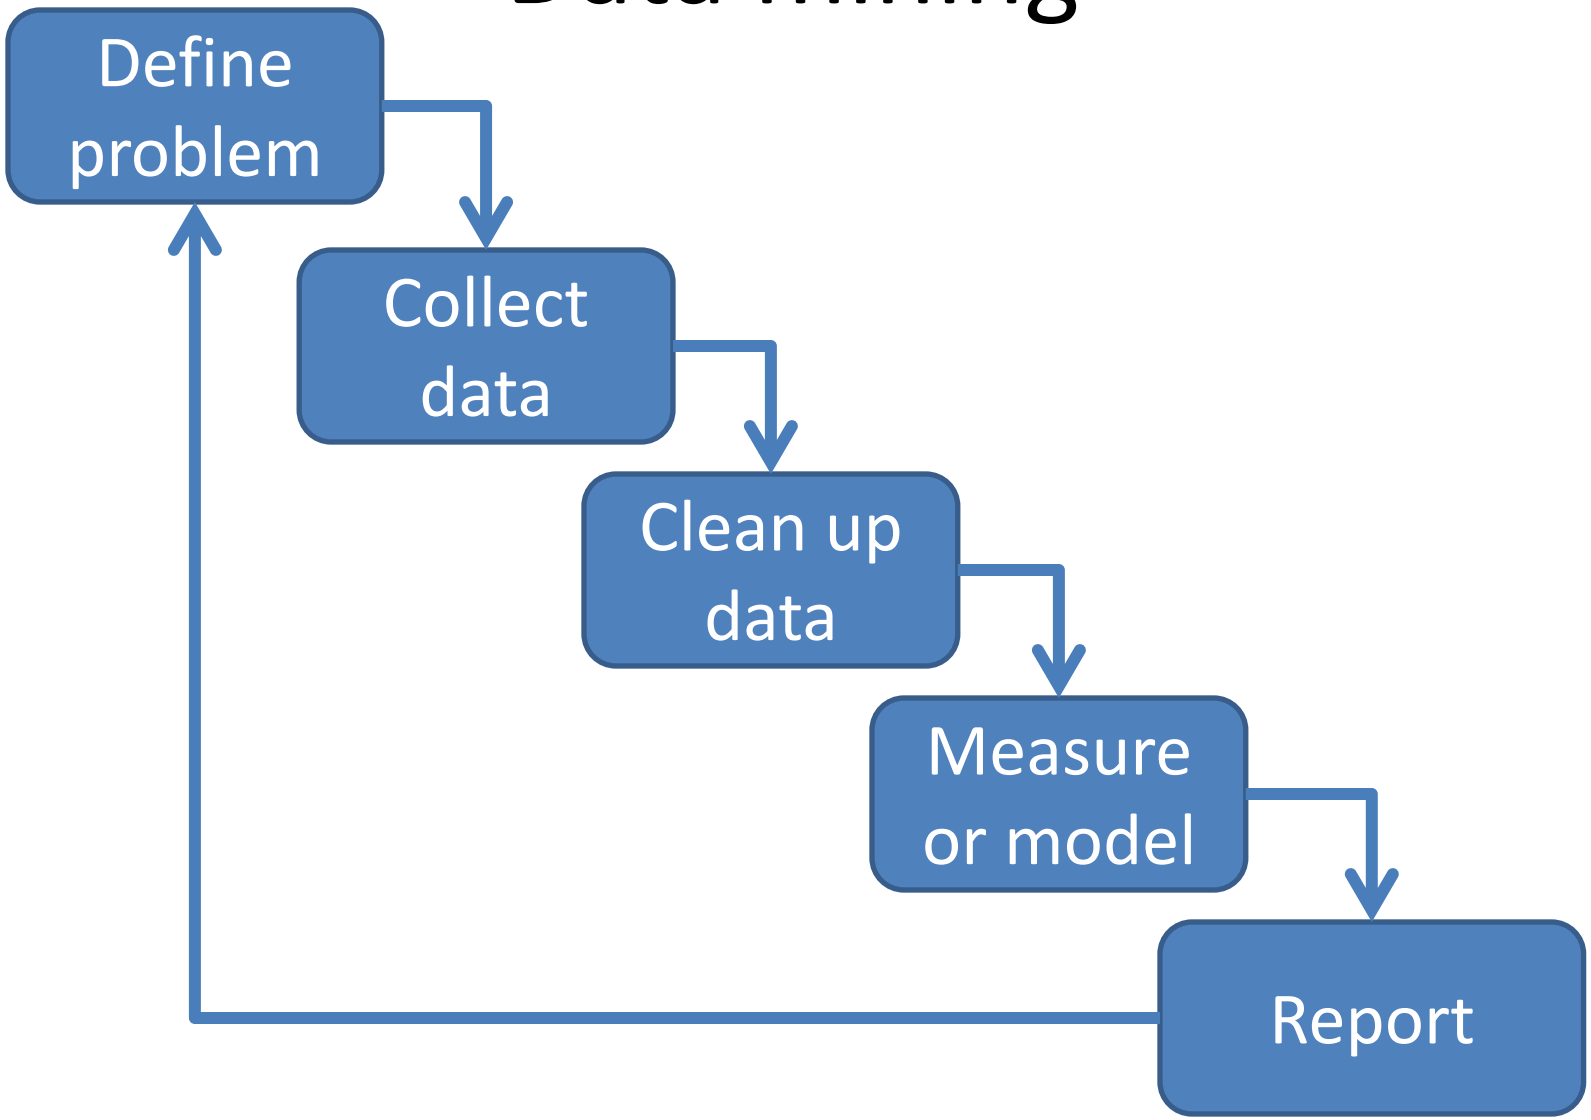
Define problem
Collect data
Clean up data
Measure or model
Report

Statistics

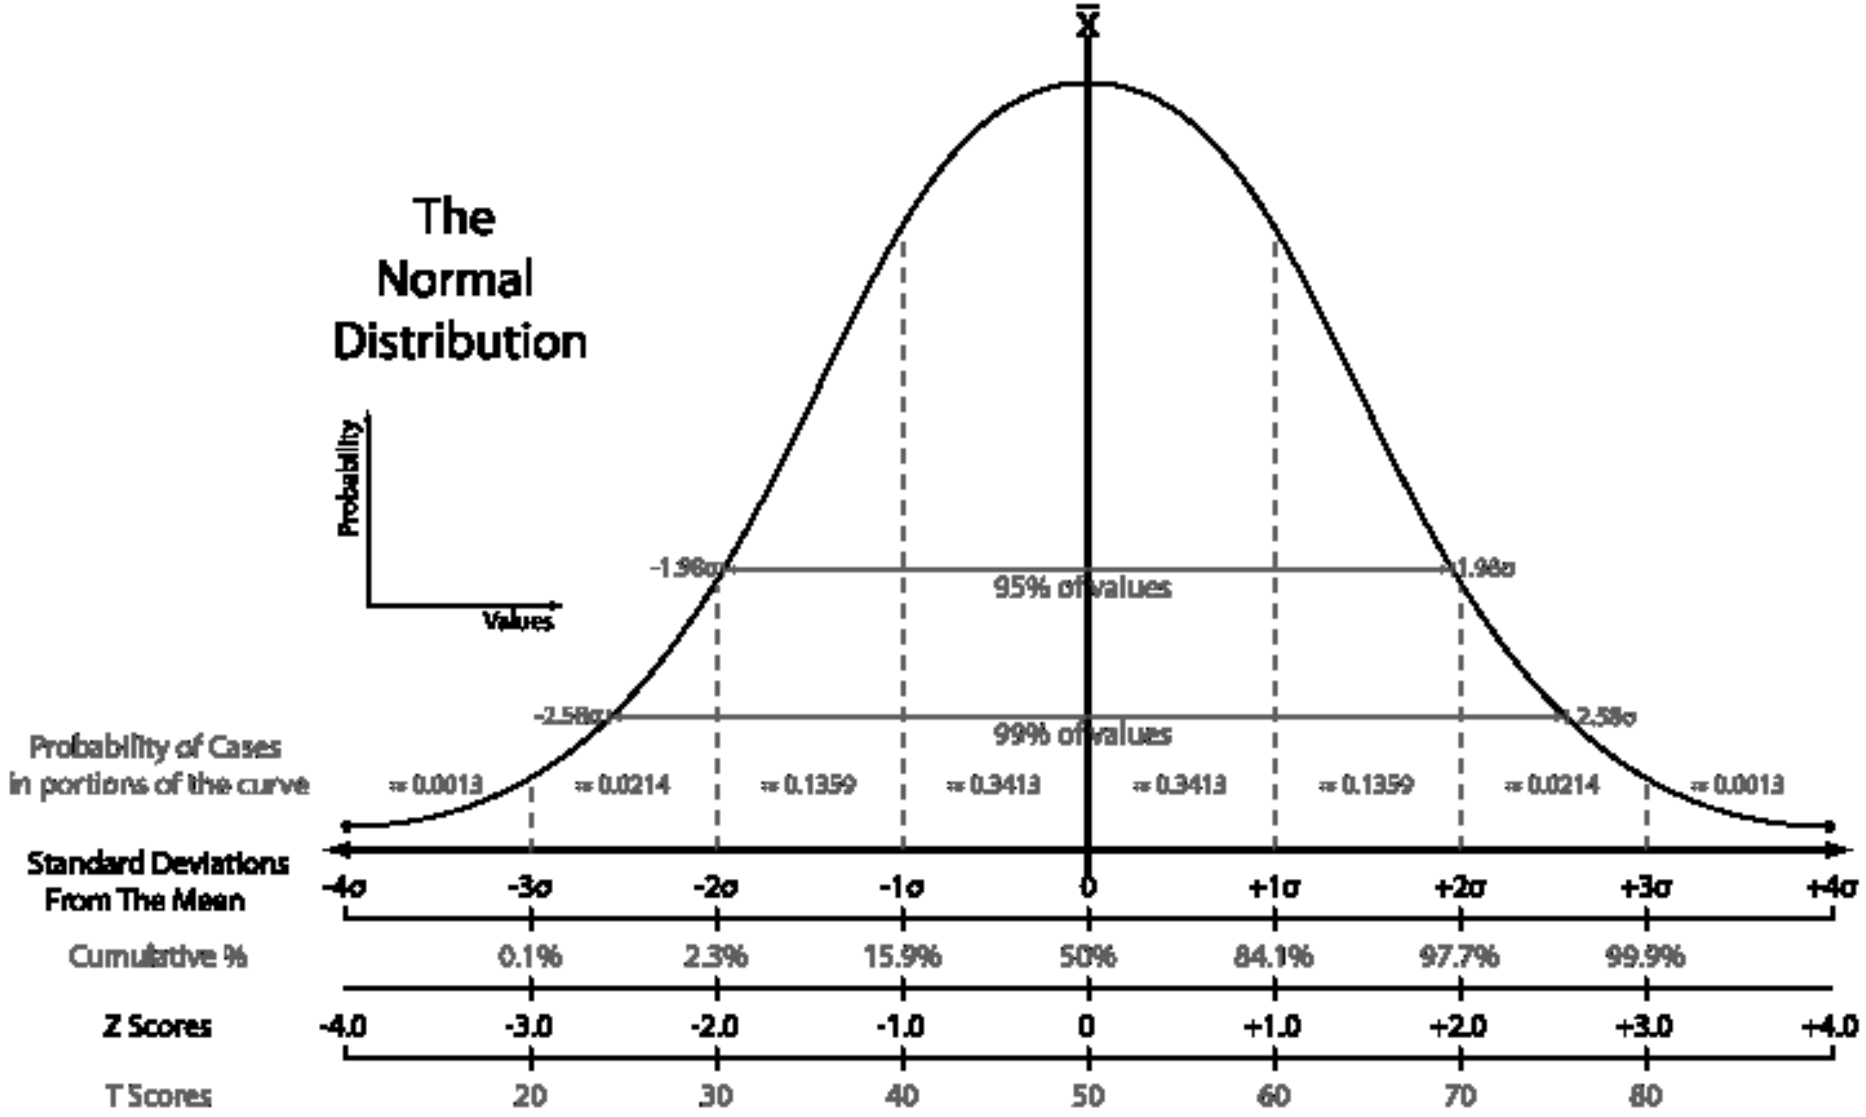
The Normal Distribution
Probability
Values
X̄
-1.96σ
+1.96σ
95% of values
-2.58σ
+2.58σ
99% of values
Probability of Cases in portions of the curve
≈ 0.0013
≈ 0.0214
≈ 0.1359
≈ 0.3413
≈ 0.3413
≈ 0.1359
≈ 0.0214
≈ 0.0013
Standard Deviations From The Mean
-4σ
-3σ
-2σ
-1σ
0
+1σ
+2σ
+3σ
+4σ
Cumulative %
0.1%
2.3%
15.9%
50%
84.1%
97.7%
99.9%
Z Scores
-4.0
-3.0
-2.0
-1.0
0
+1.0
+2.0
+3.0
+4.0
T Scores
20
30
40
50
60
70
80

Machine learning

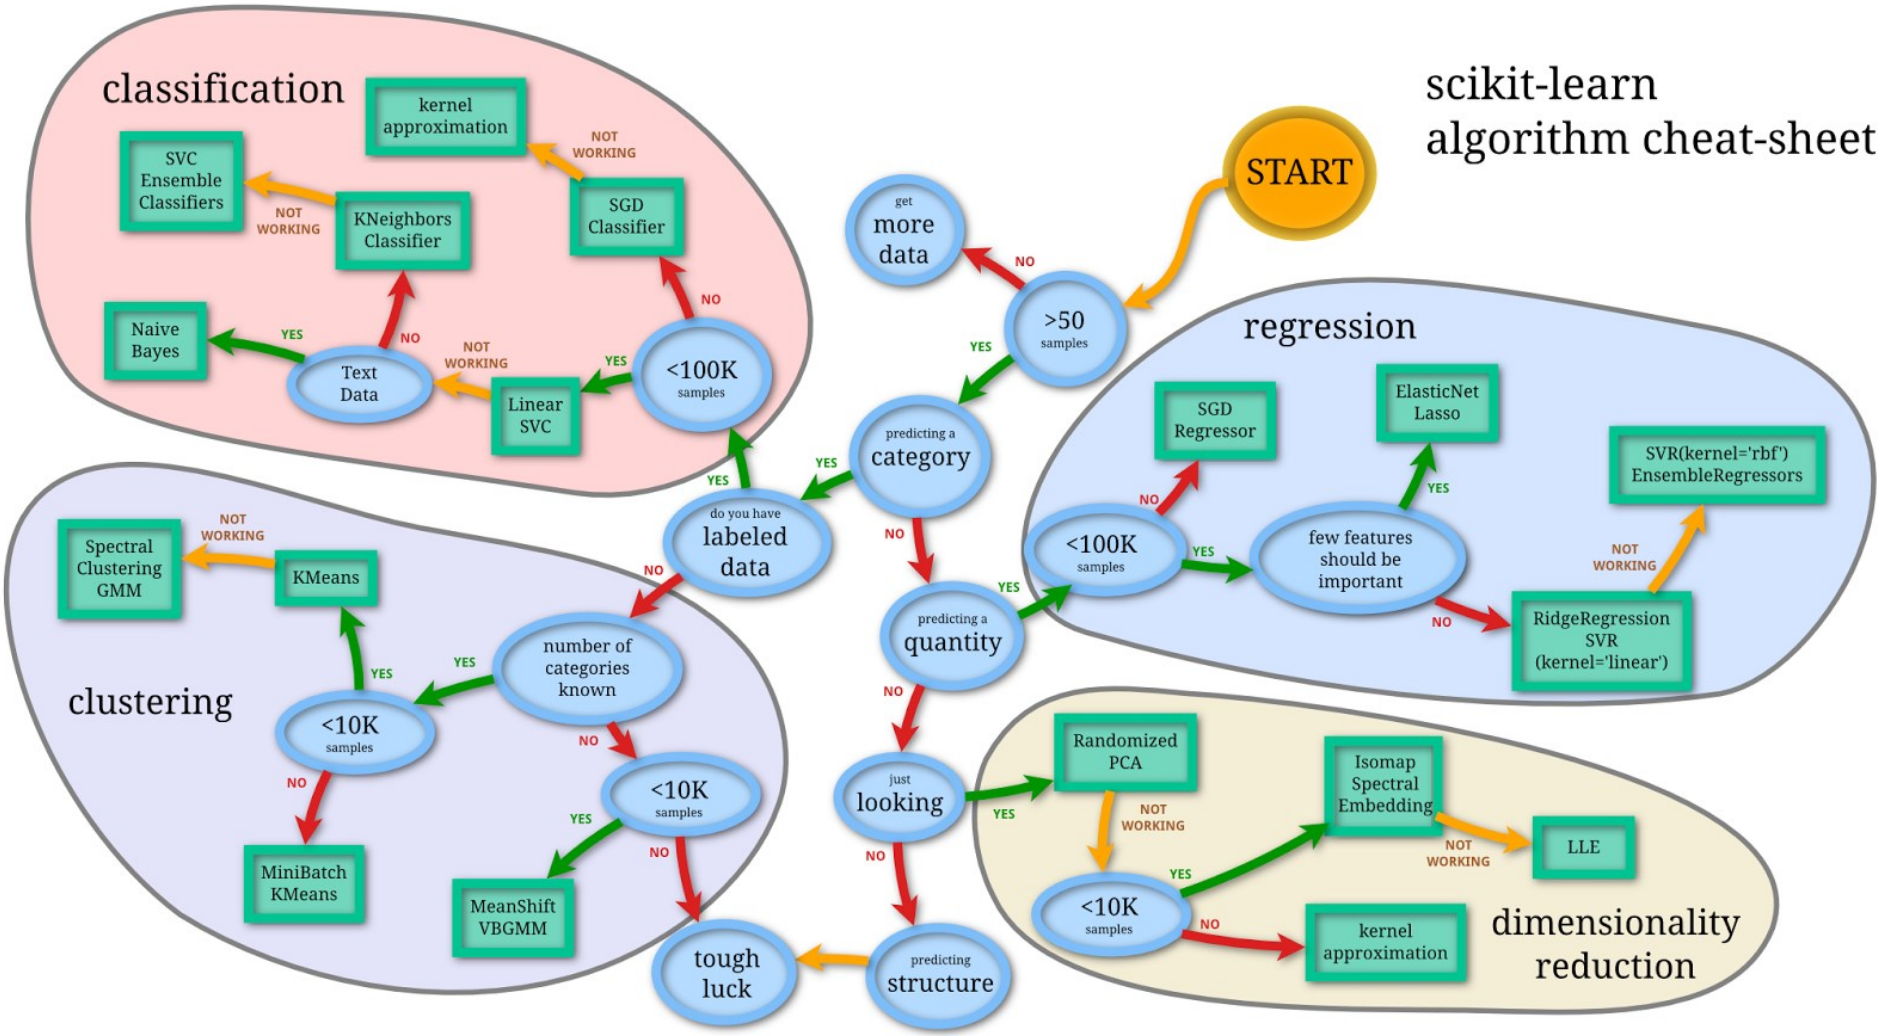

http://n-chandra.blogspot.com/2013/01/picking-machine-learning-algorithm.html

Visualization

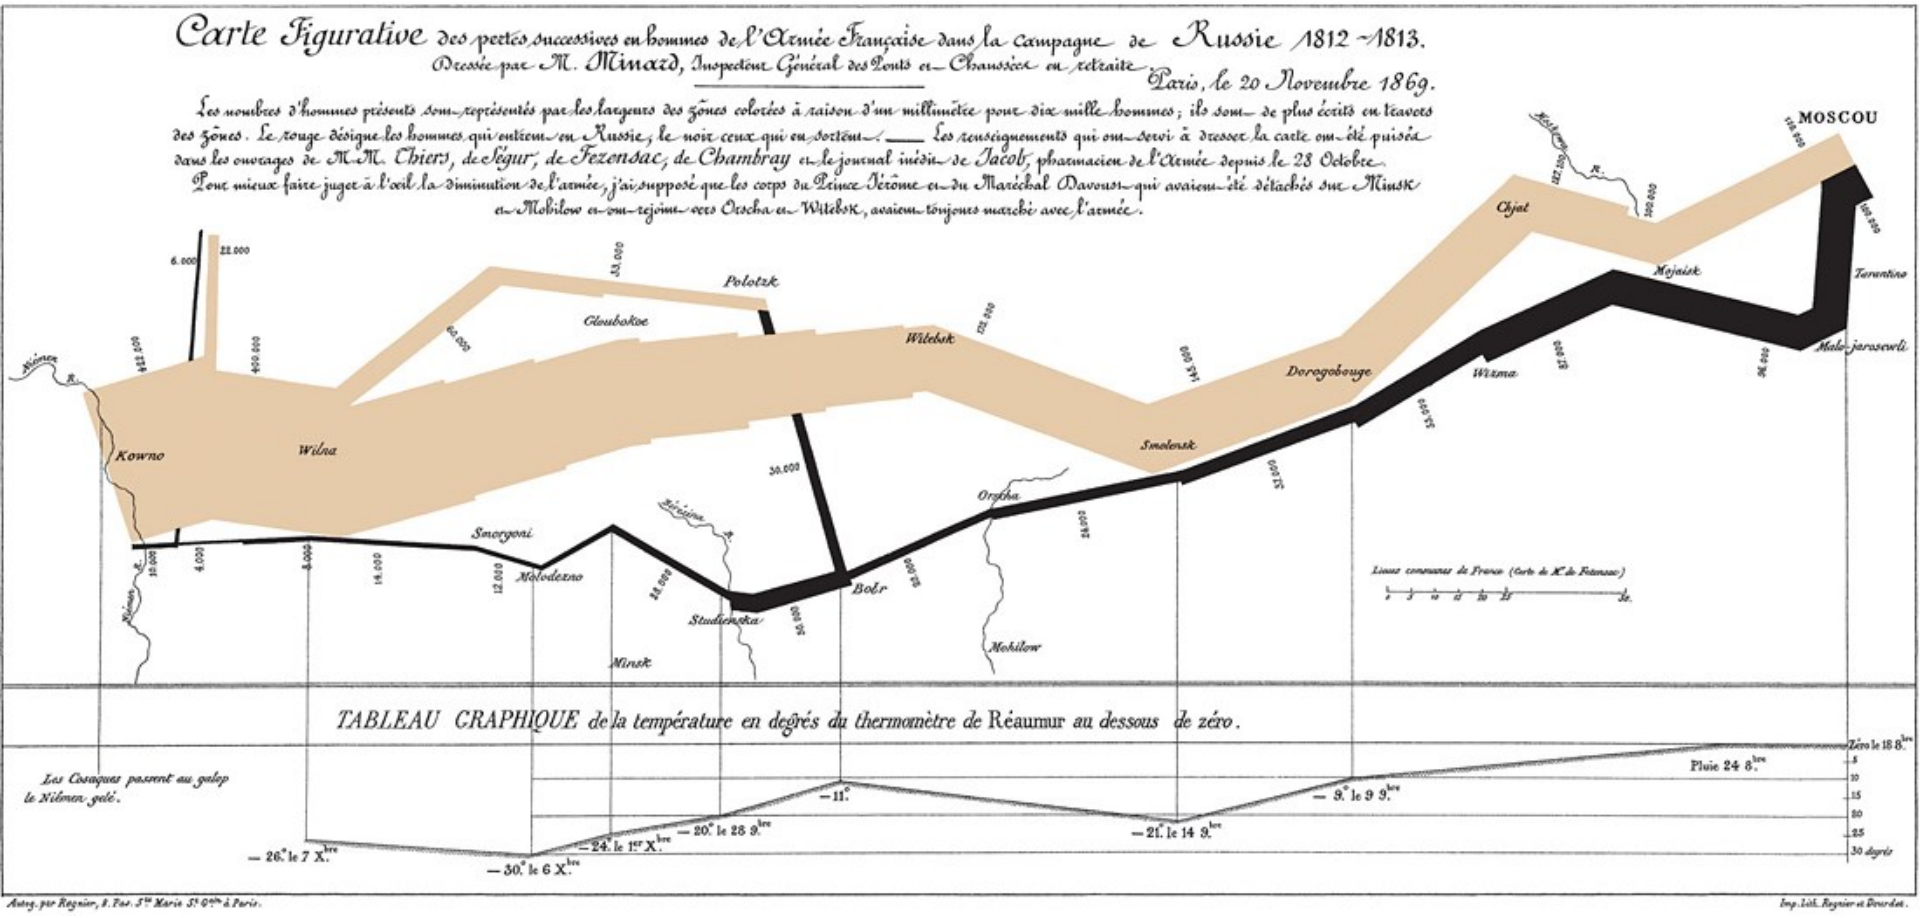

Data engineering

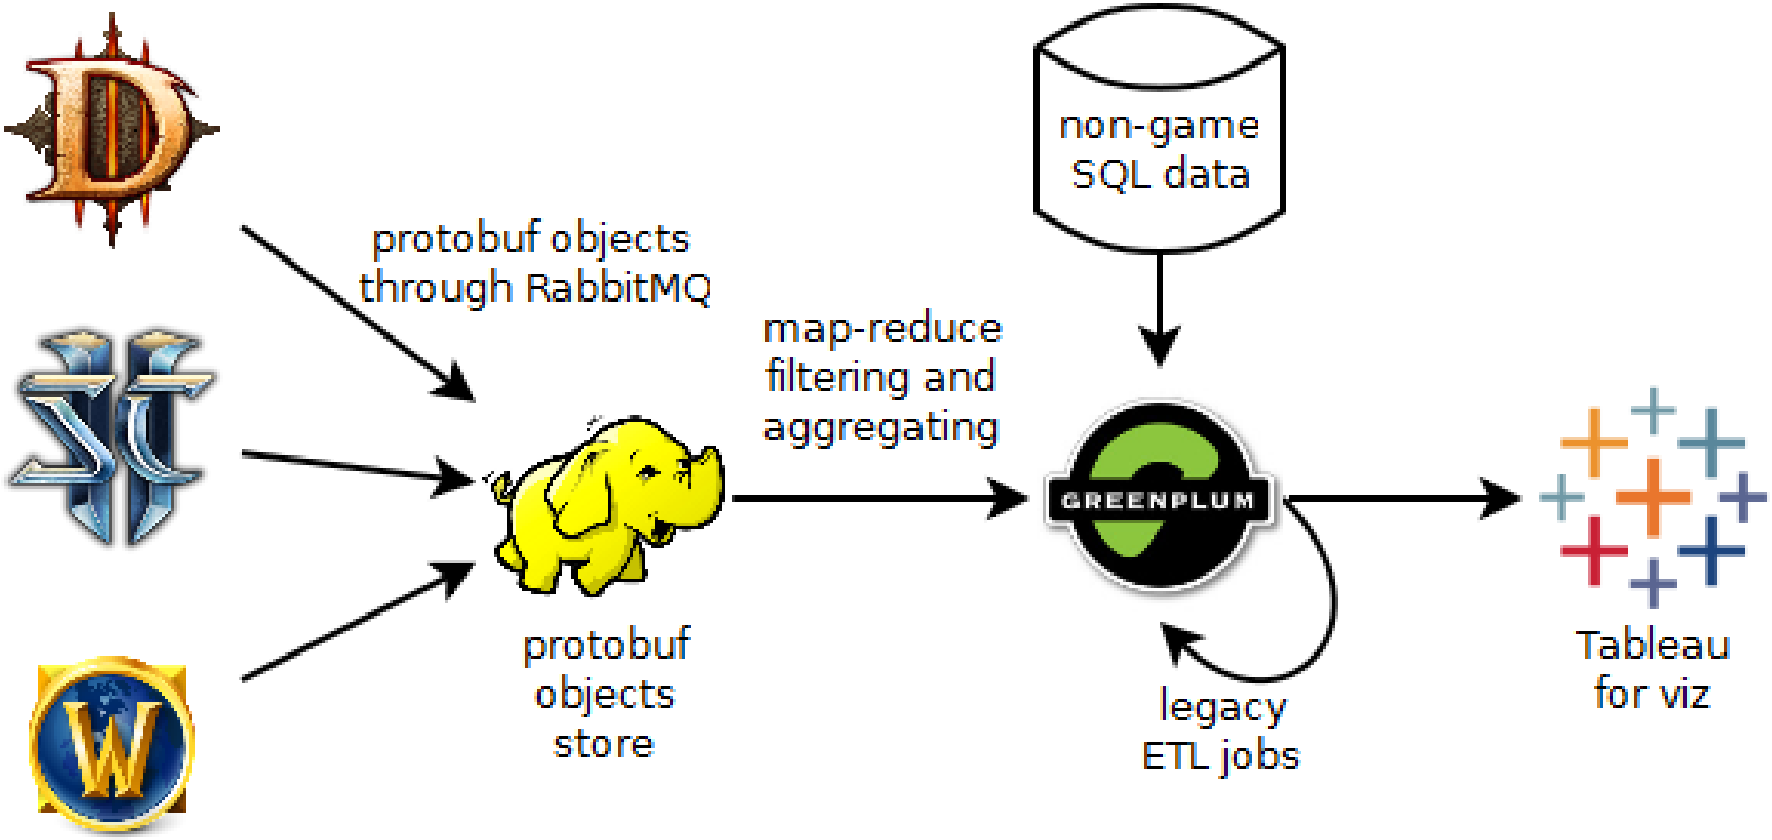

http://bowling-bash.blogspot.com/2013/10/blizzards-hadoop-platform.html

Take-aways

- Many tools and processes
- Some old, some recent
- Pick the right one for the problem at hand

INDUSTRY EXAMPLES

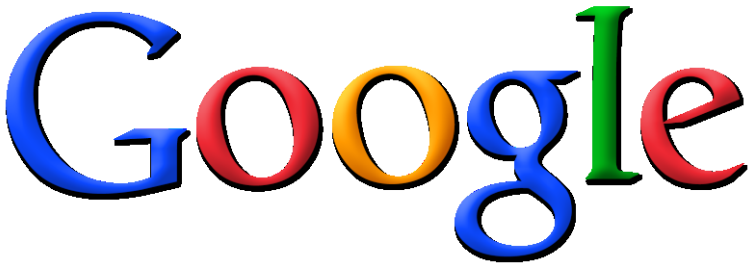
Google


trulia®
real estate search


LinkedIn

TURNSTYLE

amazon.com

NETFLIX

Take-aways

- Useful
- Everywhere
- Creepy?

SINGULAR VALUE DECOMPOSITION

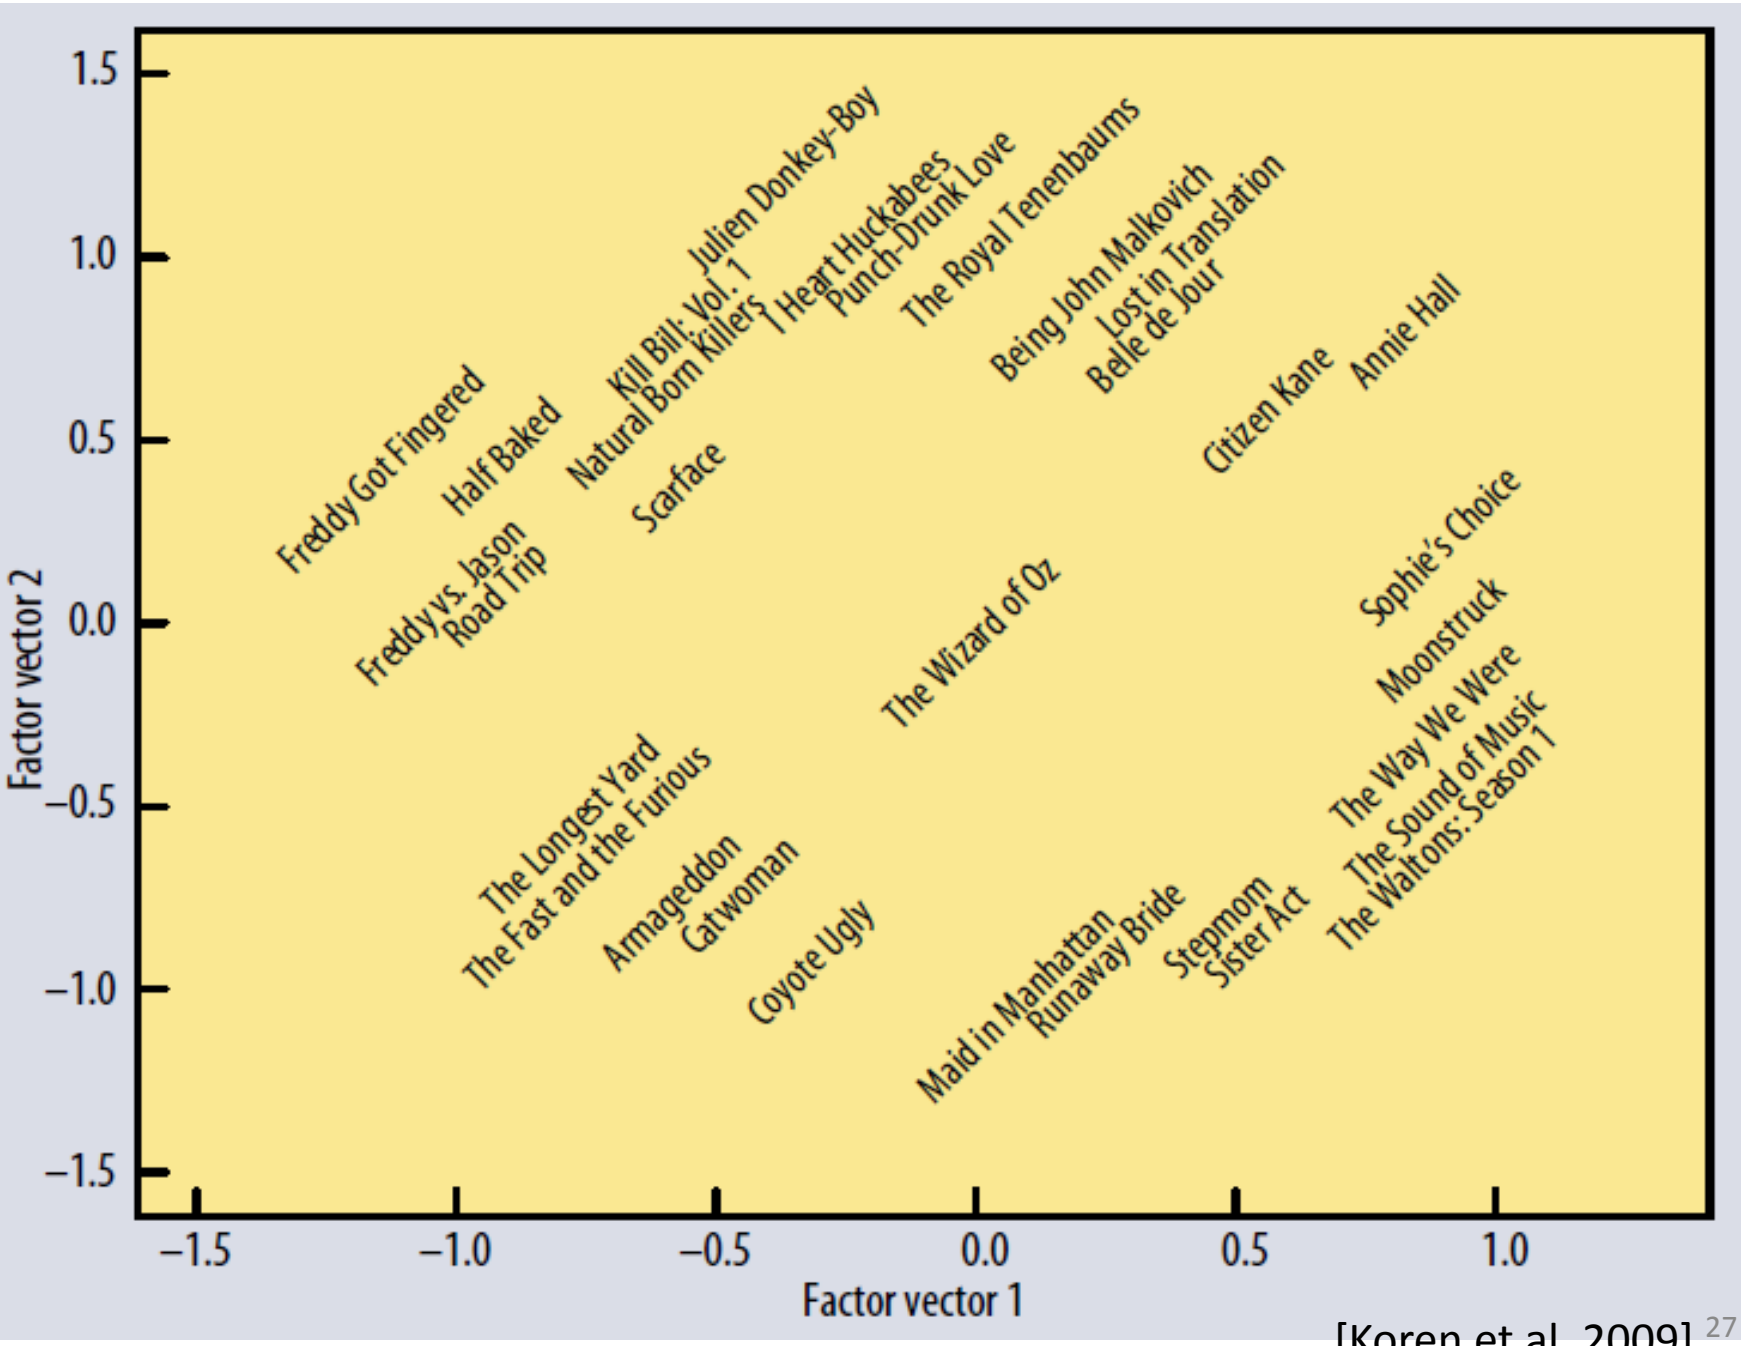

[Koren et al. 2009]

Ratings data

	The Longest Yard	Citizen Kane	Fast & Furious	Kill Bill	Sound of Music
Alice	1	4	1		5
Bob	5		4	4	
Claire		4	2		4
Dan	4	1		3	
Eve		3			5
Frank	3	3	5	2	

(usually 99% sparse)

Singular Value Decomposition

	The Longest Yard	Citizen Kane	Fast & Furious	Kill Bill	Sound of Music
Alice	1	4	1	M	5
Bob	5	M	4	4	M
Claire	M	4	2	M	4
Dan	4	1	M	3	M
Eve	M	3	M	M	5
Frank	3	3	5	2	M

=

Users

*0.1	Factor 1	Factor 2
Alice	-4	7
Bob	-5	-4
Claire	-4	2
Dan	-4	-5
Eve	-4	0
Frank	-4	0

x

Tastes

	F1	F2
F1	18	0
F2	0	4

x

Movies

*0.1	The Longest Yard	Citizen Kane	Fast & Furious	Kill Bill	Sound of Music
F1	-4	-4	-4	-4	-5
F2	-5	4	-6	1	4

M = global average = 3.3

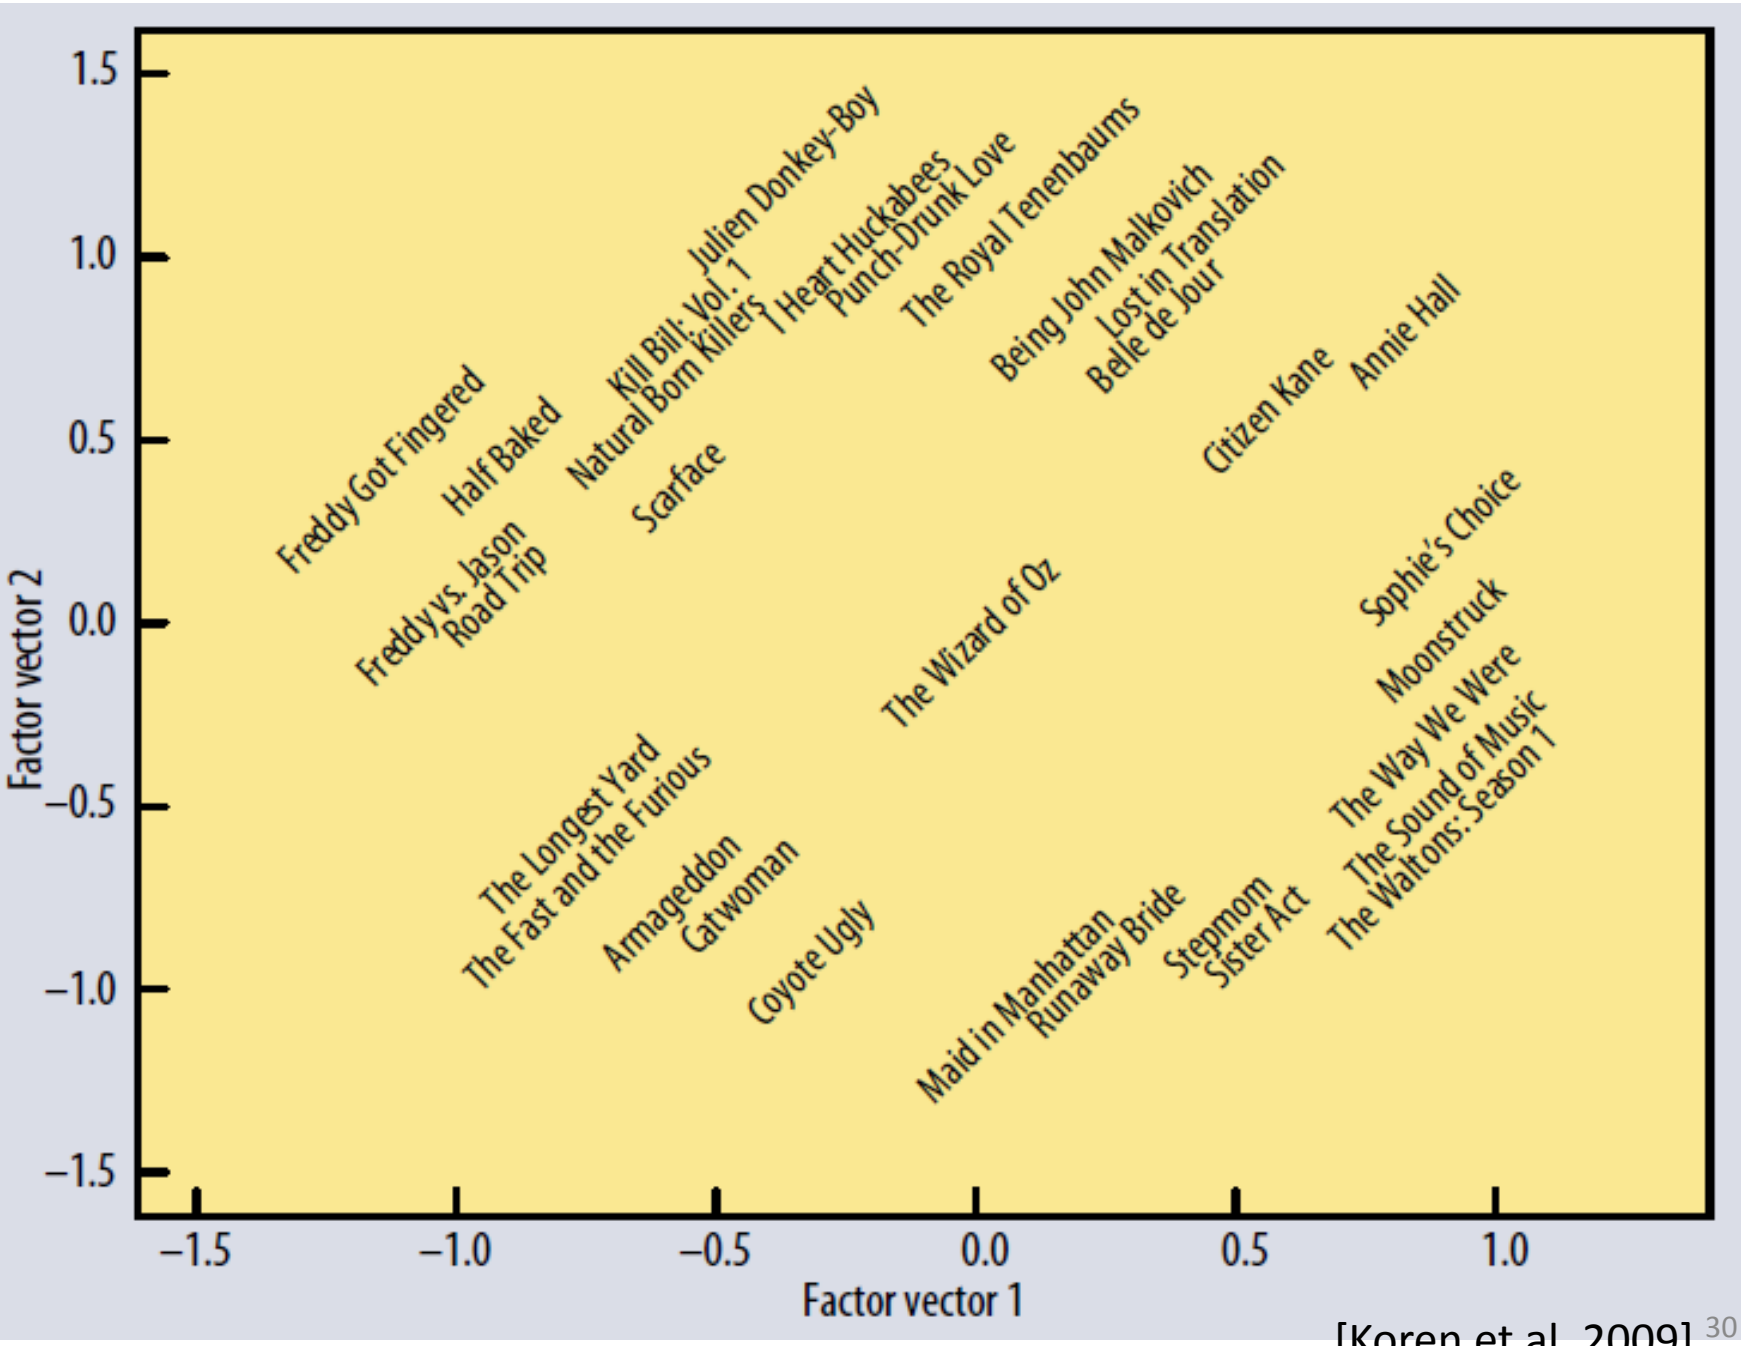

[Koren et al. 2009]

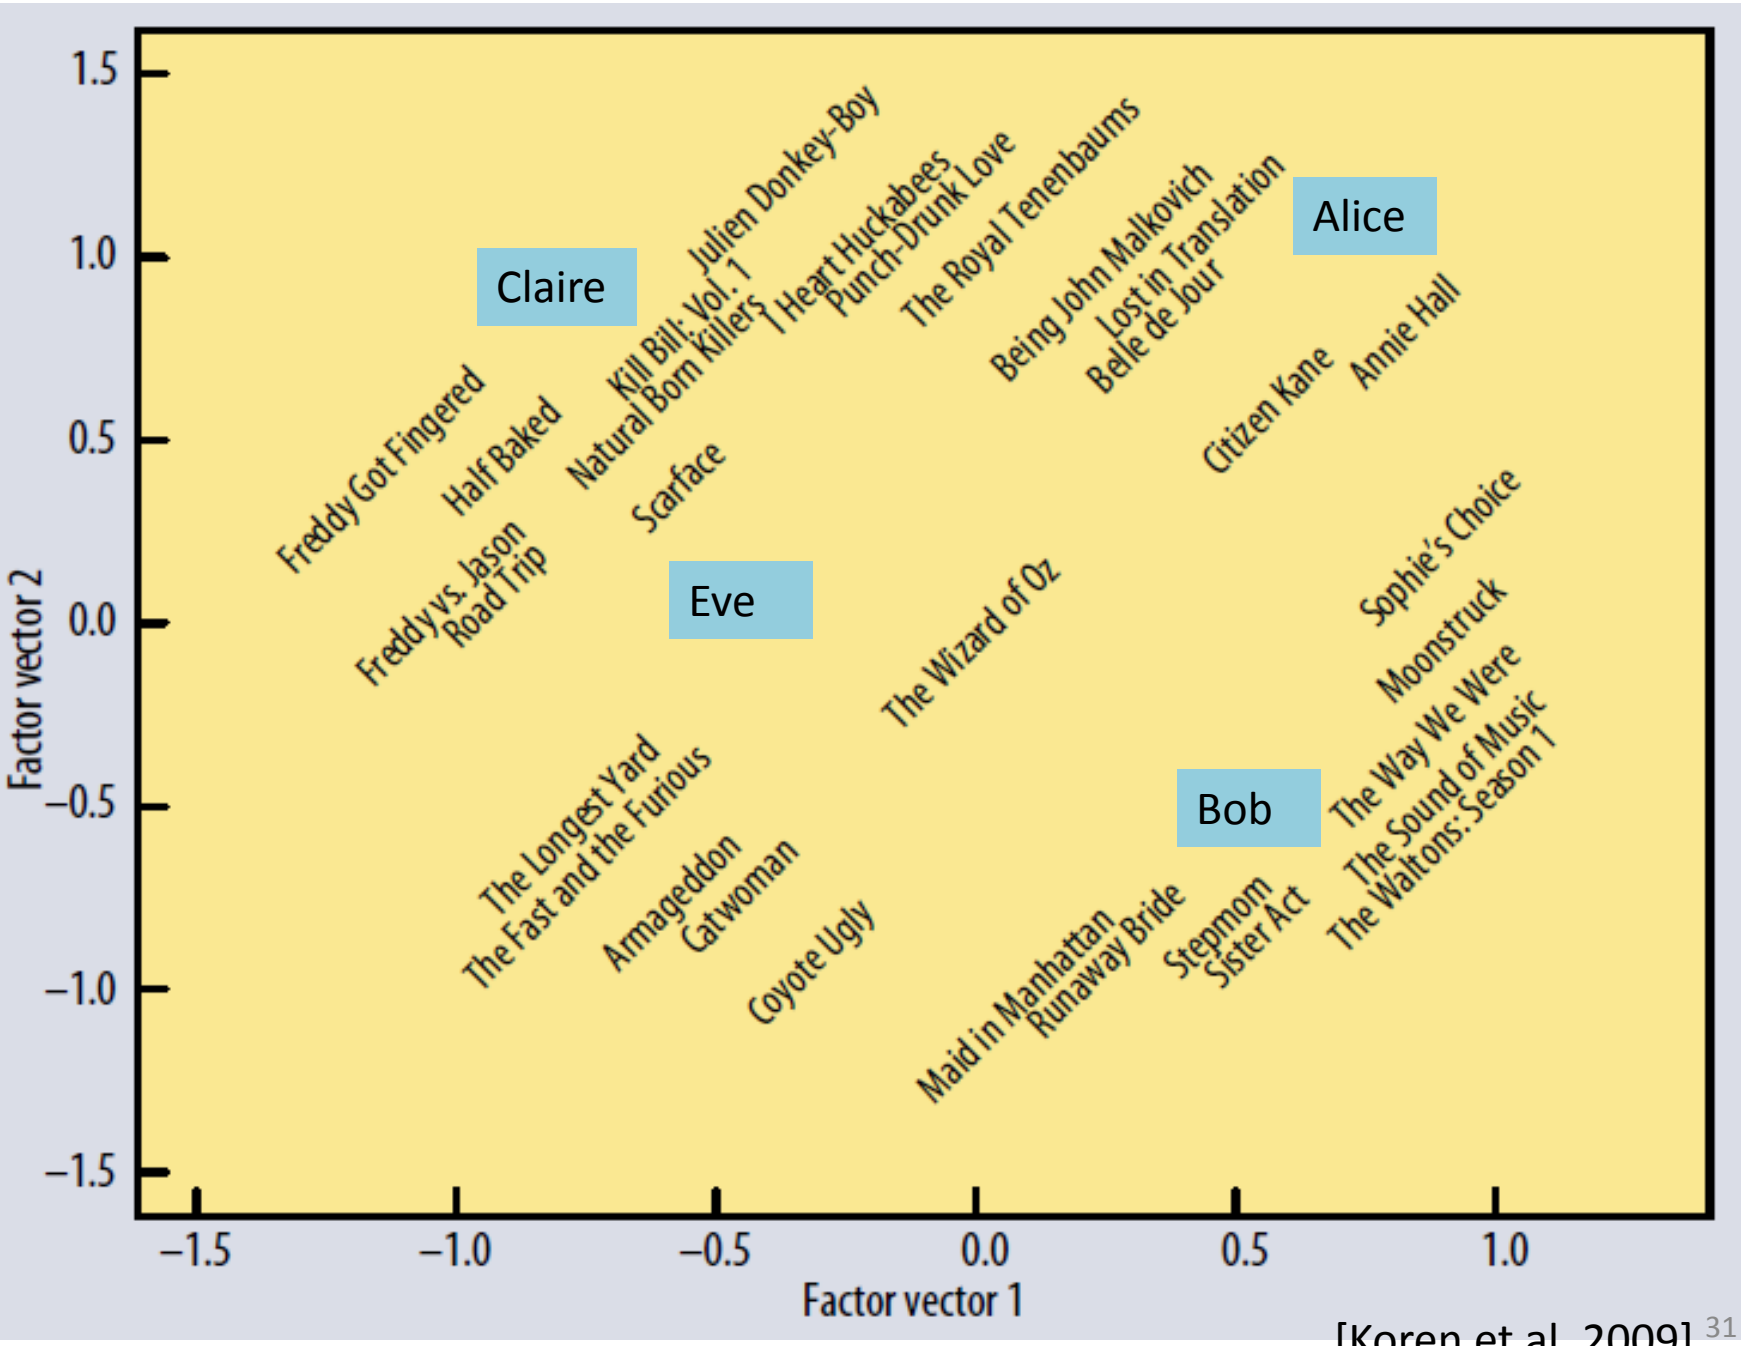

[Koren et al. 2009]

Densifying sparse matrices

Number of ratings x100

- Storage x100
- SVD is $O(N^3)$, CPU x1M

Minimize square error on known ratings using stochastic gradient descent

$$\min_{q^*,p^*} \sum_{(u,i)\in\kappa} (r_{ui} - q_i^T p_u)^2 + \lambda(|| q_i ||^2 + || p_u ||^2)$$

	Armageddon	Citizen Kane	Fast & Furious	Kill Bill	Sound of Music
Alice	1	4	1	M	5
Bob	5	M	4	4	M
Claire	M	4	2	M	4
Dan	4	1	M	3	M
Eve	M	3	M	M	5
Frank	3	3	5	2	M

Similar techniques

- Latent Semantic Analysis
  - Movies -> documents, ratings -> tf-idf
- Latent Dirichlet Allocation
  - Cf “A theory of aspects as latent topics”
- Principal Component Analysis
  - Eigen value decomposition
  - Of the covariance matrix

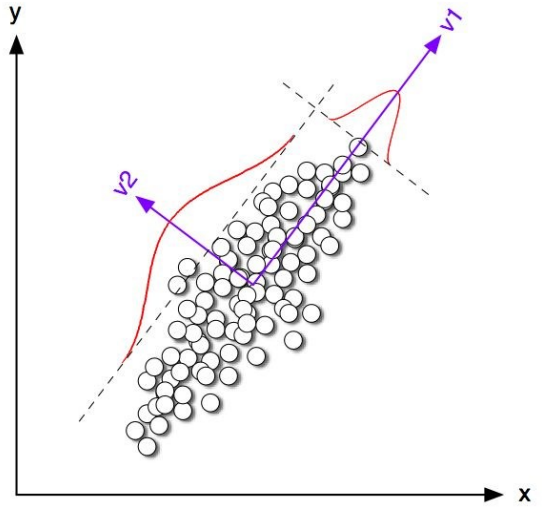

Take-aways

- Matrices
- Need practical methods
- The math escalates quickly
  - But you may need it …

HOW WOULD YOU DO …

Amazon

- “Frequently bought together”
  - Association rule mining
- Customer clustering
  - SVD

- (Stock management
  - Poisson process?)

Questions | Tags | Users | Badges | Unanswered | Ask Question

Title Is peer to peer possible with python's asyncore?

Questions that may already have your answer

2	Is is possible to code an app that can detect and connect to peers without a central servers? How?
0	is it possible peer to peer connection over wan without any server? 1
0	Python Error "Connection reset by peer" 10054
0	Web Audio API and WebRTC 1
0	can peer A connect to peer B if server is on peer B? 1

Links | Images | Styling/Headers | Lists | Blockquotes | Code | HTML | advanced help »

```
"This module provides the basic infrastructure for writing asynchronous socket service
clients and servers."
http://docs.python.org/2/library/asyncore.html#module-asyncore

If I want a non- client-server architecture, such as peer to peer, can asynchat/asyncore
still do the job?
```

draft saved

"This module provides the basic infrastructure for writing asynchronous socket service clients and servers." http://docs.python.org/2/library/asyncore.html#module-asyncore

If I want a non- client-server architecture, such as peer to peer, can asynchat/asyncore still do the job?

Tags

python | asyncore

Similar Questions

- Detecting socket close with Python's asyncore and smtpd
- asyncore python hangs
- Sending data using asyncore doesn't work
- python asyncore not keeping up with high data rates
- asyncore server: Request resulting in "socket.error'>: [Errno 32] Broken pipe)"
- Asyncore not working properly with Tkinter GUI
- Python, Asyncore and forks
- Which Python async library would be best suited for my code? Asyncore? Twisted?
- python asyncore keep track of clients
- How do I connect to IRC through a SOCKS proxy using asyncore/asynchat?
- How to re-establish asyncore connection with server (solved)
- Redirect a method call to something with a file descriptor - asyncore
- Python asyncore - multiple SMPP PDU's in one TCP packet

Dating website

- Match %
  - SVD on weighted questions
    - “Would you rather be weird or normal?”= 99% weird
    - “Sex before marriage?”
  - SVD on tf-idf from profile essays
  - Cosine distance in SVD space
- Actual information in a profile
  - Everybody loves travelling, but they say it differently
  - “I love travelling!” = “travel” = low tf-idf
  - “Have a passport!” = “passport” = high tf-idf
  - Tf-idf is not always appropriate!

Thomas Debeauvais, tdebeauv@uci.edu

THANKS

References

- Padhraic Smyth's slides on Netflix recsys http://www.ics.uci.edu/~smyth/courses/cs277/public_slides/recommender_systems_part2.pdf
- Koren et al. 2009: Matrix factorization techniques for recommender systems
- http://www.theatlantic.com/technology/archive/2014/01/how-netflix-reverse-engineered-hollywood/282679/
- http://online.wsj.com/news/articles/SB10001424052702303453004579290632128929194